

## **RISK ASSESSMENT**

### **Amazon's Proposed Acquisition of an AI-Powered Last-Mile Delivery Optimization Service**

### **Navigating Market Competition, Cybersecurity Exposure, and Cost-Reduction Opportunity in Amazon's Proposed AI-Powered Last-Mile Delivery Optimization Platform Acquisition (RIVR)**

*ALY 6130 – Risk Management Analytics*

Instructor: Abeba Nigussie Turi

Group 6: Sandra Amponsah · Truc Linh Hoang · Victoria Kpetigo

June 26, 2026

## Table of Contents

|                                                                     |    |
|---------------------------------------------------------------------|----|
| Table of Contents .....                                             | 2  |
| 1. BACKGROUND .....                                                 | 3  |
| 2. IDENTIFICATION OF RISKS .....                                    | 5  |
| 3. QUALITATIVE RISK ASSESSMENT .....                                | 10 |
| 4. QUANTITATIVE RISK ASSESSMENT .....                               | 12 |
| 5. RISK RESPONSE STRATEGY .....                                     | 18 |
| 6. Key Risk Indicators (KRIs) and Trigger Points for Action .....   | 20 |
| 7. CONCLUSION .....                                                 | 22 |
| REFERENCES .....                                                    | 24 |
| Appendix 1: Complete Risk Treatment and Response Plan (RT&RP) ..... | 26 |

## **1. BACKGROUND**

### **1.1 Company Overview and Strategic Context**

Amazon.com, Inc. is one of the world's largest players in e-commerce, cloud computing, and logistics. In its 2024 annual report, the company reported revenue increasing from \$575 billion to \$638 billion. At peak capacity, Amazon moves close to 8 million packages per day through a massive logistics network that includes more than 1,100 fulfillment, sortation, and delivery facilities, over 100 Amazon Air aircraft, tens of thousands of branded delivery vans, and more than 200,000 Amazon Flex drivers and Delivery Service Partner (DSP) associates.

Despite this scale, last-mile delivery remains Amazon's most expensive operational activity. Industry estimates place last-mile costs at \$14–\$17 per residential package, and across the sector the last mile typically represents 50–60 percent of total parcel costs. Because Amazon's routing systems were built gradually over twenty years, rather than as a unified, AI-optimized platform even small efficiency improvements can generate significant savings at Amazon's volume.

In March 2026, Amazon acquired RIVR, a Zurich-based physical-AI and robotics company spun out of ETH Zurich's Robotic Systems Lab. RIVR's platform integrates real-time data on traffic, weather, package volume, driver availability, delivery density, and past failed-delivery patterns to generate more efficient route recommendations. Strategically, the acquisition is intended to help Amazon close its optimization gap relative to emerging last-mile technologies and potentially position the company to commercialize the platform for external customers in the future.

### **1.2 The ERM Problem**

From an enterprise risk standpoint, this acquisition is far more than a routine technology upgrade. The new platform will shape core operational activities, including delivery commitments, driver instructions, customer notifications, and large-scale routing decisions, and it will process highly sensitive information such as customer addresses, real-time driver locations, proof-of-delivery images, and data from carrier APIs. If the system fails, the consequences would not be isolated. Amazon could face simultaneous operational disruptions, financial losses, regulatory exposure, and reputational harm. These risks span people, processes, technology, facilities, legal and compliance obligations, financial performance, and information security. Each area must be evaluated as part of the acquisition itself, not deferred until after integration.

At the center of the ERM challenge is the dual nature of Amazon’s proposed last-mile optimization acquisition: it introduces both significant opportunities and meaningful threats. On one hand, the platform could lower delivery costs and support Amazon’s sustainability goals. On the other hand, it brings new competitive pressures and increases exposure to cybersecurity, regulatory, and operational risks. The purpose of this report is to identify, assess, quantify, and develop responses to these risks using structured risk-management analytics.

### **1.3 How Risks Were Identified**

Risk identification began with structured brainstorming and was refined as we progressed through each course module. We started by mapping stakeholders to understand who would be affected by the acquisition: customers, drivers, DSPs, operations managers, data scientists, cybersecurity and legal teams, regulators, vendors, and even competitors. We then conducted a PESTLE scan to examine the political, economic, social, technological, legal, and environmental forces shaping AI-enabled last-mile delivery.

A pre-mortem exercise assumed the acquisition had failed two years after closing and traced backward to likely causes, including poor data quality, weak change-management practices, cybersecurity vulnerabilities, and regulatory pushback. Process mapping followed a package from the fulfillment center to the customer’s doorstep to pinpoint where AI-driven routing could introduce new risks or reduce existing ones. Finally, we translated each major risk theme into a Key Risk Indicator (KRI), creating measurable early-warning signals that leadership can monitor before, during, and after integration.

### **1.4 Competitive Context: Maturity of the Logistics-Software Market**

The competitive landscape is a critical factor because it directly shapes all three of Amazon’s priority risks. The logistics-technology software market is already well-established. UPS’s ORION system, for example, has been operating since 2012, generates an estimated \$300–\$400 million in annual savings, eliminates roughly 100 million miles each year, and required more than \$250 million and 14 years to develop (INFORMS, 2016). FedEx’s AI-driven platform, FedEx Surround, manages about 15 million shipments per day in real time (FedEx, 2024). Meanwhile, enterprise systems such as Oracle Logistics Cloud and SAP Transportation Management are deeply

embedded in tens of thousands of ERP environments, with switching costs ranging from \$5 million to \$30 million (Oracle, 2024; SAP, 2024).

By contrast, Amazon's logistics capabilities were built primarily to support internal operations rather than to compete as a commercial software provider. As a result, Amazon lacks several functions that established logistics-technology vendors rely on, such as sales engineering, SLA-based support models, customer-success teams, and competitive enterprise pricing structures. This context matters because the acquisition would push Amazon into a market environment it has not traditionally operated in. Adding to this complexity, the FTC's monopolization trial scheduled for October 2026 may further limit Amazon's pricing flexibility (Skadden, 2025).

## **2. IDENTIFICATION OF RISKS**

### **2.1 Unified Problem and Impact Statement**

Amazon's last-mile delivery network faces a set of interconnected and mutually reinforcing risks tied to the proposed acquisition of an AI-powered routing platform. These include financial overpayment, operational disruption, competitive displacement, cybersecurity breaches, algorithmic bias, low user adoption, vendor lock-in, and regulatory non-compliance. None of these exposures can be effectively managed in isolation. They require a unified enterprise-risk framework that governs every stage of the platform's acquisition, integration, and ongoing operation.

If left unaddressed, these risks would undermine the acquisition's return on investment, weaken Amazon's competitive position in both logistics operations and the logistics-software market, and increase the likelihood of enforcement actions under regulations such as the CCPA, the EU AI Act, and U.S. antitrust law. Instead of treating each risk as a standalone issue, this project frames them as components of a single, systemic exposure rooted in the acquisition decision itself.

Quantifying the exposure highlights the scale of the stakes. Failing to realize the routing-efficiency gains offered by the platform would forfeit an estimated \$400 million to \$1.36 billion in annual savings, based on a 10–25 percent per-package reduction applied to 8 million daily deliveries at \$14–\$17 per package. On the downside, regulatory penalties are substantial: CCPA violations can reach \$7,500 per intentional breach (California AG, 2023); EU AI Act fines can be as high as €35 million or 7 percent of global revenue equivalent to roughly \$50.2 billion based on Amazon's FY2025 net sales of \$716.9 billion (Financial Content, 2026); and the average cost of a data breach

is \$4.44 million globally and \$10.22 million in the United States (IBM, 2025). These figures establish a clear financial baseline: any risk-treatment strategy that costs less than the exposure it mitigates is economically justified.

## 2.2 The Consolidated Risk Register

Risk identification produced an initial set of 60 candidate risks. The register was consolidated to remove duplication between root causes and resulting risks. Opportunities were separated into their own register, and scores were recalibrated against Amazon-specific financial anchors. The final risk categories included AI strategy and governance, AI decision quality, cybersecurity and privacy, algorithmic bias and fairness, vendor and third-party dependency, financial and ROI risk, and legal and regulatory compliance. The result is a register of 32 distinct negative risks across the seven categories, scored on a Likelihood (1/3/5/7/9) by Impact (1/2/4/6/8/9) scale with a maximum score of 81. Severity tiers are Very High (score  $\geq 54$ ), High (40–53), Moderate (18–39), and Low ( $< 18$ ). Across the 32 risks the distribution is 7 Very High, 11 High, 14 Moderate, and 0 Low - a profile concentrated in the upper-middle of the matrix rather than in extreme tails.

| # | Category                        | Risks   | Count | VH/H  | Highest-severity risk (score)           |
|---|---------------------------------|---------|-------|-------|-----------------------------------------|
| 1 | AI Strategy & Governance        | R1–R4   | 4     | 1 / 1 | R1 AI strategy failure (56)             |
| 2 | AI Decision Quality             | R5–R9   | 5     | 2 / 0 | R7 Peak-ops system downtime (63)        |
| 3 | Cybersecurity & Privacy         | R10–R13 | 4     | 0 / 2 | R10 Breach of AI platform (45)          |
| 4 | Algorithmic Bias & Fairness     | R14–R17 | 4     | 0 / 2 | R14 Delivery-zone bias (45)             |
| 5 | Vendor & Third-Party Dependency | R18–R21 | 4     | 0 / 2 | R18 Vendor lock-in (42)                 |
| 6 | Financial & ROI                 | R22–R26 | 5     | 3 / 2 | R26 Cost of lagging in AI adoption (72) |

| # | Category                      | Risks   | Count | VH/H  | Highest-severity risk (score)     |
|---|-------------------------------|---------|-------|-------|-----------------------------------|
| 7 | Legal & Regulatory Compliance | R27–R32 | 6     | 1 / 2 | R27 EU AI Act non-compliance (56) |

Positive risks are kept in a separate Opportunities register and managed with Exploit or Enhance strategies. These opportunities included:

1. AI route-efficiency cost savings (9×8 = 72, High)
2. Competitive market leadership (56, High)
3. Platform commercialization (30, Medium)
4. Sustainability and ESG gains (28, Medium).




Figure 1. Risk heat map of the 32 consolidated negative risks plotted on the Likelihood (1/3/5/7/9) by Impact (1/2/4/6/8/9) matrix. Cells are colored by severity tier: Very High (≥ 54), High (40–53), Moderate (18–39), and Low (< 18). The distribution — 7 Very High, 11 High, 14 Moderate, and 0 Low — concentrates in the upper-middle of the matrix. R7 (peak-operations downtime) and R26

*(cost of lagging in AI adoption) sit at the highest-severity cells. The four opportunities are managed in a separate register (Table above) and are not plotted here.*

### **2.3 The Three Priority Risks Carried Forward**

Three priority risks were selected for deeper quantitative analysis because they span the competitive, cybersecurity, financial, and operational dimensions of the proposed acquisition, each has an observable leading indicator. They are labeled PR1–PR3 here to avoid confusion with the register numbering (R1–R3 in the register refer to AI-strategy risks); the appended I&W notebook labels them R1–R3 internally. The table maps each to its canonical register or opportunity entry.

| <b>Priority risk</b> | <b>Description</b>                                    | <b>Type</b> | <b>Canonical cross-reference</b>                                                                                                                     |
|----------------------|-------------------------------------------------------|-------------|------------------------------------------------------------------------------------------------------------------------------------------------------|
| <b>PR1</b>           | Commercial market-entry / enterprise win-rate failure | Threat      | From the SWOT Threats quadrant; relates to R26 (cost of lagging in AI adoption) and the downside of opportunities O2/O3. Not a single register line. |
| <b>PR2</b>           | Cybersecurity breach of geolocation and address data  | Threat      | R10 Cybersecurity breach of AI routing platform (45, High) and R11 Customer/driver data-privacy breach (40, High).                                   |
| <b>PR3</b>           | Cost-reduction and sustainability leadership          | Opportunity | Opportunity O1 AI route-efficiency cost savings (72, High).                                                                                          |

#### ***PR1 - Entry into the competitive logistics software market***

Amazon’s acquired platform will enter a market dominated by long-established competitors such as UPS ORION, FedEx Surround, Oracle Logistics Cloud, and SAP Transportation Management, vendors with decades of experience, mature customer relationships, and deep carrier integrations. The challenge for Amazon is structural rather than technical. Its logistics division was built to support internal operations, not to function as a commercial software business, meaning it lacks



essential capabilities like sales engineering, enterprise-grade SLA support, customer-success teams, and competitive pricing models.

This structural gap creates a predictable risk pathway: incumbents respond with aggressive pricing and rapid feature releases, Amazon's enterprise win-rate drops below 40 percent, and commercial revenue falls short of expectations, ultimately weakening the acquisition's ROI. Amazon's strongest differentiator, the training data generated from 8 million daily stops, cannot be monetized externally without a fully developed enterprise sales function. With a Likelihood of 7 and an Impact of 6, this risk scores 42 (High).

**Key Risk Indicator:** enterprise win-rate and market-share growth

- **Green:** win-rate above 50% and market-share growth above 8% annually
- **Amber:** win-rate between 40–50% or growth between 3–8%
- **Red:** win-rate below 40% for two consecutive quarters or declining market share

### ***PR2 - Cybersecurity and data-breach exposure***

The platform will handle highly sensitive logistics data, including customer home addresses, real-time driver geolocation, proof-of-delivery photos, route intelligence, and carrier API feeds across a distributed architecture. A ransomware attack, unauthorized API access, or major breach would disrupt delivery operations, trigger CCPA multi-state notification within 72 hours, expose Amazon to regulatory enforcement, and erode customer trust—especially given that Amazon has not previously faced this level of scrutiny over logistics-data governance.

The risk is most acute during the 18–24-month integration period, before zero-trust controls, hardened APIs, and full monitoring capabilities are fully deployed. A breach would turn Amazon's richest asset—the density of its delivery data, into a multi-jurisdictional liability. With a Likelihood of 5 and an Impact of 9, this risk scores 45 (High), defined by its asymmetry: moderate probability but extremely high regulatory and reputational consequences.

**Key Risk Indicator:** high-severity vulnerabilities and unauthorized-access events

- **Green:** no critical vulnerabilities past SLA; no anomalous access
- **Amber:** critical vulnerability unresolved for 7–15 days or anomalous API traffic detected
- **Red:** critical vulnerability unresolved beyond 15 days or any confirmed unauthorized access

### ***PR3 - Cost-reduction and sustainability leadership (opportunity)***

AI-enabled route optimization can significantly reduce miles driven, fuel consumption, idle time, and manual dispatching. Industry research estimates a 15–25% per-package cost reduction compared with rule-based routing (Shuaibu et al., 2025). At Amazon’s scale, even a 10% improvement translates into hundreds of millions of dollars in annual savings. These efficiency gains also reduce emissions, helping Amazon stay ahead of California’s Clean Miles Standard and EU urban-access regulations, turning compliance requirements into a competitive advantage.

This opportunity is time-sensitive: UPS ORION’s benchmark compounds each year, and regulatory windows will not remain open indefinitely. Realizing the upside depends heavily on whether drivers and DSPs consistently follow AI-generated route recommendations. Scored as a positive risk at Likelihood 9 and Impact 8, this opportunity carries a 72 (High upside).

**Key Risk Indicator:** per-package delivery-cost reduction relative to baseline

**Green:** >15% reduction

**Amber:** 5–15% reduction

**Red:** <5% reduction, with driver/DSP adoption tracked as the leading sub-indicator

### 3. QUALITATIVE RISK ASSESSMENT

#### 3.1 Scoring Methodology

Scores were recalibrated under three principles. A likelihood of 9 (at least 80 percent probability) is reserved for events that are near-certain given current conditions, such as capital costs or the cost of inaction. Root causes and resulting risks are merged into a single entry rather than counted twice. Impact scores are anchored to financial ranges, with Impact 9 (over \$5 billion) reserved for existential exposures such as an EU AI Act penalty. The probability assignments are documented expert judgments grounded in cited evidence, not externally sourced probability figures, which keeps the scoring transparent and defensible.

#### 3.2 Analytical Framework: Scenario Analytics and Industry Fusion Analytics

Two analytical methods from Fleisher and Bensoussan (2015) were applied to the three priority risks. Scenario Analytics models each risk across optimistic, base-case, and pessimistic futures based on its key uncertainty drivers. This pushes the analysis beyond single-point estimates and shows how each risk evolves over time. Industry Fusion Analytics then benchmarks each risk

against documented precedents in adjacent industries, grounding probability and impact estimates in real evidence. Used together, scenario analysis highlighted the time sensitivity of each risk, while fusion analysis provided the external reference points that informed the quantitative modeling.

### ***PR1 - Commercial market-entry***

The main uncertainty drivers include the speed at which Amazon can build a sales-engineering function, how aggressively incumbents respond to pricing and features, the strength of enterprise lock-in, and antitrust constraints on Amazon's conduct. In the base case, Amazon enters the market without a dedicated sales function, and its enterprise win-rate stabilizes around 38 percent for three quarters, enough to generate internal ROI but insufficient to meet external commercialization targets.

Fusion benchmarks reinforce this vulnerability window. The Oracle–PeopleSoft integration required roughly 36 months to mature its sales motion, and Google Cloud's entry against AWS and Azure saw win-rates below 30 percent for three years despite strong technology. Both cases show that enterprise buyers prioritize readiness, SLAs, support, and compliance documentation over product quality alone. Because the win-rate is observable quarterly, this risk is well-suited for Indicators & Warnings (I&W) tracking.

### ***PR2 - Cybersecurity breach***

Key drivers include the pace of zero-trust deployment, the sophistication of threat actors, the enforcement posture of the California AG and FTC, and architectural choices during integration. In the base case, zero-trust controls take about nine months to reach all 1,100+ stations, creating a partial-control window in months 6–9. During this period, an attempted attack is detected through anomalous geo-API traffic before any confirmed data exfiltration.

Fusion benchmarks underscore the severity of the downside tail: the IBM/Ponemon (2025) U.S. average breach cost of \$10.22 million, CCJ Digital's (2026) finding that 33 percent of logistics-sector AWS credentials have been exposed, and the 2021 Colonial Pipeline attack, which caused more than \$4 billion in cascading disruption. This asymmetric profile, with moderate probability but extreme regulatory and operational consequences, is exactly the type of risk where

EMV and Monte Carlo tail modeling provide the most value. As a result, PR2 moves into quantification with the highest urgency.

### ***PR3 - Cost-reduction opportunity***

The main drivers are driver and DSP adoption of AI-generated routes, training-data quality for under-represented delivery zones, regulatory timelines under the Clean Miles Standard and EU urban-access rules, and competitor response speed. In the base case, adoption reaches roughly 70 percent within twelve months—below the 80 percent target, and per-package cost falls by about 14 percent, generating an estimated \$490 million in annual savings.

Fusion benchmarks validate the upside potential: UPS ORION’s \$300–\$400 million annual savings, 100 million miles removed, and 100,000 tonnes of CO<sub>2</sub> avoided (INFORMS, 2016); Walmart’s 30 percent AI-driven shipping-cost reduction (2024); and the 15–25 percent efficiency range reported by Shuaibu et al. (2025). Because the per-package cost reduction is directly measurable and compounds over time, this opportunity qualifies for quantification and should be actively pursued rather than simply monitored.

All three priority risks, therefore, advance to quantitative analysis. **PR2** presents the most severe downside tail, **PR1** carries the largest internally controllable expected loss, and **PR3** offers the greatest positive expected value.

## **4. QUANTITATIVE RISK ASSESSMENT**

The figures in this section come directly from the team’s reproducible notebook (Group6\_IW\_Quantitative\_Risk\_Assessment.ipynb, SEED = 42, 50,000 Monte Carlo trials). Every real-world dollar value is anchored to an official source in the benchmark registry; the indicator panel is synthetic and seeded because the target’s internal KPIs are not public, with magnitudes calibrated to those benchmarks. This satisfies the requirement to use a realistic proxy dataset with stated assumptions where real data is unavailable.

### **4.1 Methodology: Indicators and Warning (I&W)**

The I&W technique (Fleisher & Bensoussan, 2015) turns observable leading indicators into a calibrated probability. For each priority risk, we defined four indicators anchored to the established KRIs, weighted them to sum to 1.0, normalized each to a 0–1 warning contribution, and summed

them into a composite warning score over a 156-week (36-month) horizon; a composite of 0.55 or above flags a red warning state. The event probability blends the indicator-derived estimate with the Module-3 register anchor as  $P = 0.7 \times (\text{warning-derived}) + 0.3 \times (\text{register anchor})$ . For the positive risk PR3, the warning-derived term is  $(1 - \text{warning})$ , since a low warning means the opportunity is being realized.

**Threshold calibration.** In direct response to the Module-3 feedback that warning thresholds appeared generic, every threshold below is anchored to a published benchmark in the logistics, enterprise-sales, or cybersecurity domain rather than set by convention, so that a red state corresponds to a level of exposure external evidence already treats as unacceptable. The 0.55 composite red trigger illustrates the logic: on the Module-3 non-linear scale (maximum Likelihood  $\times$  Impact = 81), the High band opens at a score of 40 ( $40/81 \approx 0.49$ ) and the Very-High band at 54 ( $54/81 \approx 0.67$ ), so a composite of 0.55 fires the alarm at the point the qualitative register already classifies as requiring active executive mitigation, consistent with the firm’s risk appetite under ISO 31000:2018 and COSO ERM (2017). The per-indicator thresholds are each tied to an external benchmark, as summarized in the table below.

**Table 1. Warning-threshold calibration and empirical/theoretical basis**

| Threshold                                                             | Where used                                           | Benchmark basis (why this cut-off)                                                                                                                                                                                                                                         | Source                            |
|-----------------------------------------------------------------------|------------------------------------------------------|----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|-----------------------------------|
| Composite $\geq 0.55 \rightarrow$ Red                                 | Master escalation trigger (all three priority risks) | Maps to the High-to-Very-High band of the Module-3 L $\times$ I register (High = $40/81 \approx 0.49$ ; Very High = $54/81 \approx 0.67$ ), so the alarm fires at the exposure the firm already treats as needing active mitigation.                                       | ISO 31000:2018; COSO ERM (2017)   |
| PR1: win-rate $< 40\%$ for two consecutive quarters $\rightarrow$ Red | win-rate-gap indicator                               | Competitive B2B win-rates average $\approx 47\%$ (RAIN Group) and $\approx 21\%$ blended (HubSpot, 2024). Against entrenched incumbents (ORION, Surround) a rate held below 40% means losing most head-to-head evaluations; two quarters filters quarter-to-quarter noise. | RAIN Group (2023); HubSpot (2024) |
| PR2: critical vulnerability open $> 15$ days $\rightarrow$ Red        | vulns-past-SLA indicator                             | CISA BOD 19-02 sets 15 days as the remediation deadline for critical vulnerabilities on internet-facing systems (30 for high), now the de-facto private-                                                                                                                   | CISA BOD 19-02; IBM (2025)        |

| Threshold                                                            | Where used               | Benchmark basis (why this cut-off)                                                                                                                                                                                                              | Source                                      |
|----------------------------------------------------------------------|--------------------------|-------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|---------------------------------------------|
|                                                                      |                          | sector SLA; reinforced by IBM/Ponemon's 241-day mean breach lifecycle.                                                                                                                                                                          |                                             |
| PR3: per-package cost reduction < 5% → opportunity not materializing | cost-reduction indicator | Documented AI route-optimization savings run 15–30%; the conservative INFORMS/ORION floor is $\approx$ 5% (ORION delivered 10–20% over manual routing). Below 5%, the platform underperforms the weakest benchmark, and the value thesis fails. | INFORMS/ORION (2016); Shuaibu et al. (2025) |

## 4.2 Probability Distributions, Expected Loss, Variance, and Exposure

Probability was modeled with a Beta distribution and conditional impact with a three-point PERT distribution, where

**Expected Impact** =  $(\text{Optimistic} + 4 \times \text{Most Likely} + \text{Pessimistic}) / 6$

**Expected Monetary Value (EMV)** =  $\text{Probability}(\text{event}) \times \text{Expected (Impact)}$

Impact variance is  $((P - O) / 6)^2$ , and exposure is summarized by the 95th-percentile (P95) outcome from the simulation. The blended probabilities reconcile closely with the qualitative register scores, which validates the earlier judgment while putting a defensible number on it.

| Risk | P(event) | O / M / P (\$M)      | E[Impact] (\$M) | EMV (\$M) | Var (\$M <sup>2</sup> ) | P95 exposure (\$M)      |
|------|----------|----------------------|-----------------|-----------|-------------------------|-------------------------|
| PR1  | 0.65     | 200 / 500 / 1,200    | 567             | −368      | 27,778                  | 890 downside (VaR)      |
| PR2  | 0.47     | 250 / 450 / 900      | 492             | −231      | 17,361                  | 701 downside (VaR)      |
| PR3  | 0.62     | +400 / +700 / +1,200 | +733            | +453      | 44,444                  | +991 upside (best case) |

PR3 shows the widest variance (\$44,444 M<sup>2</sup>), reflecting the spread between a stalled pilot (\$400M) and a fully scaled deployment (\$1,200M). The gap between EMV and P95 is largest for PR1 (−\$368M against \$890M), which is the practical signal that contingency reserves set at the mean alone would be under-provisioned; reserves should be sized against the P95 figure.

#### 4.3 The I&W Indicator Panel

| Risk       | Leading indicators (diagnostic weights)                                                                                     | Mean warning | Blended P(event)                                   |
|------------|-----------------------------------------------------------------------------------------------------------------------------|--------------|----------------------------------------------------|
| <b>PR1</b> | win-rate gap (.40), competitor releases (.20), sales-engineering vacancy (.20), benchmark deficit (.20); red at $\geq 0.55$ | 0.61         | $0.7 \times 0.61 + 0.3 \times 0.743 \approx 0.65$  |
| <b>PR2</b> | critical vulns past 15-day SLA (.35), anomalous geo-API (.25), zero-trust gap (.25), MTTR days (.15)                        | 0.44         | $0.7 \times 0.44 + 0.3 \times 0.540 \approx 0.47$  |
| <b>PR3</b> | per-package cost reduction (.40), driver adoption (.25), OTDPI (.20), pilot on-target (.15)                                 | 0.44         | (1–warn) blended w/ anchor<br>$0.748 \approx 0.62$ |

#### 4.4 Monte Carlo Simulation (50,000 Trials)

Each of the 50,000 trials draws a probability realization from the Beta distribution and a conditional-impact realization from the PERT distribution, producing a full exposure distribution for each risk rather than a single point. The mean gives expected impact and the 95th percentile gives the tail figures reported above (PR1 about \$890M downside, PR2 about \$701M downside, PR3 about \$991M upside). It is the width of these distributions, not just their averages, that justifies contingency reserves and risk-appetite thresholds, because PR1 and PR2 both carry fat downside tails that a point estimate would hide.




Figure 2. Monte Carlo simulation of risk exposure (50,000 trials). Panels are labeled R1–R3 as in the appended notebook (PR1–PR3 in this report). The gold line marks expected impact and the dashed line marks the 95th-percentile exposure - downside for PR1/PR2, upside for PR3.

#### 4.5 Machine-Learning Early-Warning Detection

To confirm that each warning state is detectable in advance, Gradient-Boosting and Logistic-Regression classifiers were trained on the indicator panel and validated with 5-fold stratified cross-validation, using ROC-AUC as the score and feature importance to identify the dominant signal. The feature set for each model is the same set of indicators that the corresponding KRI monitors, so the model relies on signals management can actually track rather than on latent statistical constructs.

| Risk       | ROC-AUC (GBM) | Dominant early-warning signal     | Interpretation and lead time                                                                                                                                                              |
|------------|---------------|-----------------------------------|-------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| <b>PR1</b> | 0.87          | sales-engineering vacancy gap     | Unfilled sales-engineering roles precede win-rate decline by about two to three months; without technical sales coverage Amazon cannot win enterprise RFPs regardless of product quality. |
| <b>PR2</b> | 0.89          | Anomalous geolocation-API traffic | Unauthorized geo-API access precedes a confirmed breach by roughly 15 to 45 days - the highest-value early-warning signal in the portfolio.                                               |



| <b>Risk</b> | <b>ROC-AUC<br/>(GBM)</b> | <b>Dominant early-warning<br/>signal</b> | <b>Interpretation and lead time</b>                                                                                                                |
|-------------|--------------------------|------------------------------------------|----------------------------------------------------------------------------------------------------------------------------------------------------|
| <b>PR3</b>  | 1.00 (Logistic)          | driver adoption of AI routes             | On the synthetic panel, adoption is a near-deterministic gate on savings realization, which makes it the single most actionable lever in the deal. |

#### 4.6 Integrated Situation Assessment

The model’s inputs are tied directly to specific business and non-business conditions and making those links explicit shows how the external environment shapes the distributions presented above. On the business side, Amazon’s \$14–\$17 per-package delivery cost anchors the impact distribution for PR3: a \$400 million gain corresponds to roughly a 10 percent reduction applied to 8 million daily packages at \$14 each. The maturity of the logistics-software market, entrenched incumbents, \$5M–\$30M switching costs, and deep Oracle/SAP ERP lock-in compresses PR1’s optimistic scenario below what Amazon’s product quality alone might imply. Likewise, the 33 percent exposure rate of logistics-sector AWS credentials (CCJ Digital, 2026) is what sets the 0.25 weight on PR2’s zero-trust-gap indicator.

On the non-business side, the EU AI Act’s fine of up to 7 percent of global revenue (roughly \$50.2 billion for Amazon) acts as a regulatory floor that widens PR2’s pessimistic tail and forces investment in zero-trust controls and cyber-insurance regardless of probability estimates. California’s Clean Miles Standard creates a time-bounded window for PR3’s \$1,200 million optimistic case, since the upside only materializes if Amazon achieves early compliance. The FTC antitrust trial scheduled for October 2026 further narrows PR1’s optimistic distribution by constraining pricing flexibility. Finally, the documented targeting of logistics infrastructure most notably the Colonial Pipeline attack supports holding PR2’s integration-window probability near 0.50, higher than Amazon’s historical breach rate. Excluding these inputs would understate PR2’s tail by roughly \$200 million and overstate PR1’s best-case scenario by ignoring antitrust-driven pricing limits. For this reason, the situation assessment is part of the model’s calibration, not background commentary.

## 5. RISK RESPONSE STRATEGY

Risk responses follow the four standard treatment strategies used throughout the register—**Mitigate**, **Transfer**, **Exploit/Enhance** (for positive risks), and **Accept**—consistent with ISO 31000:2018. The strategy for each risk was selected by evaluating it against three baselines (the severity threshold, Amazon’s risk appetite, and the expected loss) and then resolving four questions in sequence:

1. Is the risk above appetite;
2. Can the underlying cause be avoided cost-effectively;
3. Can the exposure be transferred or shared;
4. Can the remaining residual risk be accepted under monitoring?

The table summarizes the resulting decisions, with detailed explanations provided in the sections that follow.

| <b>Risk</b> | <b>Strategy</b>        | <b>Residual EMV</b> | <b>KRI red trigger</b>                                                                                    |
|-------------|------------------------|---------------------|-----------------------------------------------------------------------------------------------------------|
| <b>PR1</b>  | Mitigate               | −\$239M<br>(−35%)   | Win-rate < 40% for two consecutive quarters, or sales-engineering headcount < 75%                         |
| <b>PR2</b>  | Mitigate +<br>Transfer | −\$104M<br>(−55%)   | Any confirmed unauthorized geo/address API access; any critical vuln > 15 days; zero-trust coverage < 90% |
| <b>PR3</b>  | Exploit + Enhance      | +\$534M<br>(+18%)   | Per-package reduction < 5%; driver adoption < 60% (or < 80% after 6 months); OTDPI < 97 for two months    |

### 5.1 PR1 – Mitigate

PR1 stems from a structural issue: Amazon lacks a commercial software sales function, and no amount of product refinement can remove that root cause. Mitigation is still viable because Amazon’s proprietary delivery-data advantage makes the platform commercially valuable once the right sales infrastructure is in place. Since enterprise buyers prioritize readiness, SLAs, support, compliance documentation over pure product quality, the mitigation steps must be intentional and front-loaded.

Amazon should begin by commissioning an independent benchmark comparing its platform against UPS ORION and FedEx Surround before any commercial launch, then close the most critical capability gaps within twelve months. It should also hire 15–20 enterprise sales engineers with Oracle or SAP TM backgrounds before responding to the first RFP, as the ML model shows the sales-engineering gap provides two to three months of early warning before win-rates decline. Tiered introductory pricing, designed within FTC conduct limits, will help counteract the \$5M–\$30M switching costs that protect incumbents. Finally, Amazon should monitor the win-rate dashboard from day one and pause external commercialization if win-rate remains below 40 percent for two consecutive quarters, shifting focus back to internal optimization. Taken together, these actions reduce PR1’s expected monetary value (EMV) from –\$368 million to approximately –\$239 million, a 35 percent improvement.

## **5.2 PR2 – Mitigate and Transfer**

PR2’s asymmetric profile moderate probability but extremely high downside, requires a dual strategy: reduce the likelihood of a breach and cap the remaining tail risk. Amazon’s scale allows it to secure logistics-specific cyber-insurance at rates smaller carriers cannot access, while the integration timeline makes mitigation urgent because the elevated-exposure window closes once zero-trust deployment is complete.

Zero-trust should be treated as a hard launch gate: no production customer or driver data should flow through the platform until controls are fully deployed, eliminating the lateral movement that turns anomalies into confirmed breaches. A dedicated API security gateway should monitor geolocation, address, and proof-of-delivery calls for behavioral and volumetric anomalies, terminating suspicious sessions and escalating to the CISO within fifteen minutes directly addressing the 33 percent credential-exposure finding.

Cyber-insurance transfers the residual tail, covering CCPA multi-state notification, forensics, regulatory defense, and business interruption up to \$50 million per incident. A logistics-specific incident-response playbook—distinct from Amazon’s e-commerce procedures and tested through tabletop exercises before go-live—ensures the organization can meet the 72-hour CCPA notification requirement and coordinate DSP communications.

Combined, these measures reduce PR2's EMV from –\$231 million to roughly –\$104 million, the largest proportional reduction in the portfolio, with zero-trust lowering probability and insurance capping the impact tail.

### **5.3 PR3 – Exploit and Enhance**

PR3 represents a positive risk, so the strategy is to increase both the likelihood and magnitude of the upside rather than simply observe it. Amazon's dense delivery data and AWS compute capacity make the projected savings achievable at a speed competitors cannot match. The binding constraint is adoption: the ML model identifies driver and DSP adherence to AI-generated routes as the strongest predictor of realized savings.

The rollout should expand from the three-city pilot to more than 50 stations within eighteen months, prioritizing high-density markets where routing complexity is greatest and human dispatchers have the least advantage. A structured adoption program—role-specific training, in-cab coaching, a documented feedback loop, and incentives tied to AI-route acceptance—should target 80 percent adoption within six months of each launch, the threshold between amber and green on the KRI.

Early compliance with California's Clean Miles Standard converts a regulatory requirement into a competitive moat, while a public CO<sub>2</sub> KPI benchmarked against ORION's 100-million-mile reduction strengthens ESG positioning and reinforces internal momentum. Modeled together, these actions increase PR3's EMV from +\$453 million to approximately +\$534 million, an 18 percent uplift.

## **6. Key Risk Indicators (KRIs) and Trigger Points for Action**

Each KRI was built through the six CRISP-DM phases. Business Understanding defines why the metric matters to the risk; Data Understanding identifies the source systems and their quality; Data Preparation normalizes the raw reading to a 0–1 warning contribution; Modeling sets the indicator's weight in the composite warning score; Evaluation confirms predictive validity through the ML model; and Deployment fixes the monitoring frequency and escalation path. Every KRI is designed to give 30 to 90 days of advance warning, the window in which correction costs far less than the event itself, and each draw on the dominant early-warning signal identified in Section 4.5.

| Risk       | KRI (dominant signal)                                                       | Trigger thresholds                                                                                                                                                                                                                                   | Action on Red                                                                                                     |
|------------|-----------------------------------------------------------------------------|------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|-------------------------------------------------------------------------------------------------------------------|
| <b>PR1</b> | Enterprise win-rate vs ORION/Surround; sales-engineering headcount          | <b>Green:</b> win-rate $\geq 50\%$ and headcount $\geq 90\%$ .<br><b>Amber:</b> win-rate 40–49% or headcount 75–89%.<br><b>Red:</b> win-rate $< 40\%$ for two quarters, or headcount $< 75\%$ .                                                      | Pause external commercialization; VP review; add 5 sales engineers within 60 days; revisit pricing.               |
| <b>PR2</b> | Anomalous geo-API activity; critical-vulnerability SLA; zero-trust coverage | <b>Green:</b> no vulns past SLA, no anomalies, 100% coverage.<br><b>Amber:</b> 1–2 vulns near SLA or coverage 90–99%.<br><b>Red:</b> any confirmed unauthorized access; any critical vuln $> 15$ days; coverage $< 90\%$ .                           | Immediate API lockdown; start CCPA 72-hour clock; activate IR playbook; forensics within 4 hours; engage insurer. |
| <b>PR3</b> | Per-package cost reduction; driver adoption; OTDPI                          | <b>Green:</b> reduction $\geq 15\%$ , adoption $\geq 80\%$ , OTDPI $\geq 100$ .<br><b>Amber:</b> reduction 5–14% or adoption 60–79%.<br><b>Red:</b> reduction $< 5\%$ ; adoption $< 60\%$ (or $< 80\%$ after 6 months); OTDPI $< 97$ for two months. | Suspend market expansion; escalate driver training and incentives; review model drift; escalate to COO.           |

Threshold rationale. Every trigger point is anchored to an external benchmark rather than set generically. The 40-percent win-rate floor sits well below the  $\approx 47$ -percent average competitive B2B win-rate reported by RAIN Group (and the  $\approx 21$ -percent blended figure in HubSpot, 2024), so a rate held under 40 percent for two quarters means Amazon is losing most head-to-head enterprise evaluations against entrenched incumbents; the Google Cloud fusion benchmark, where sustained sub-30-percent win-rates signaled structural rather than cyclical failure, corroborates the floor and gives roughly a year of warning. The 15-day critical-vulnerability SLA follows CISA BOD 19-02, which sets 15 days as the federal remediation deadline for critical vulnerabilities on

internet-facing systems and now functions as the de-facto private-sector standard; IBM/Ponemon's 241-day average breach lifecycle confirms this keeps exposures well inside the window in which automated exploit tooling spreads. The 0.55 composite red state maps to the High-to-Very-High band of the qualitative register ( $40/81 \approx 0.49$ ;  $54/81 \approx 0.67$ ), keeping the quantitative trigger consistent with the firm's risk appetite under ISO 31000:2018 and COSO ERM (2017). The 5-percent cost-reduction floor is the conservative lower bound of documented AI route-optimization savings (an INFORMS/ORION floor of  $\approx 5\%$  against a 15–30% industry range), below which the opportunity is not materializing; the 80-percent adoption target reflects the modeling finding that savings flatten beneath it. The adoption KRI also separates overrides by reason code, so a safety-justified override is not counted the same as one driven by distrust, which keeps the signal pointed at the right management response.

## 7. CONCLUSION

Amazon's proposed acquisition is strategically sound and financially justified—but only under specific conditions. The value does not come from the technology alone; it depends on managing three interconnected risks: entering a mature and highly competitive software market, exposing a new and sensitive cybersecurity attack surface, and capturing an efficiency opportunity whose value declines with delay. The qualitative analysis validated all three risks for quantification and provided the external benchmarks that anchor the model. The quantitative analysis then translated those risks into dollar values and measurable indicators.

PR2 carries the most severe downside tail, with a P95 loss of \$701 million despite a moderate probability of 0.47. This is why zero-trust must be a non-negotiable launch requirement and why the remaining tail risk should be transferred to an insurer. PR1 shows a  $-\$368$  million expected loss, and its earliest warning signal is the sales-engineering gap—meaning Amazon must build that capability before commercial launch rather than reacting after win-rates fall. PR3 offers a  $+\$453$  million expected gain (P95  $+\$991$  million), compounding by roughly \$400 million per year relative to ORION's benchmark. This makes driver adoption the primary lever for realizing the upside. After applying the recommended treatments, the portfolio nets to approximately  $+\$191$  million ( $-\$239$ M for PR1,  $-\$104$ M for PR2,  $+\$534$ M for PR3), not including the strategic option value of owning the platform.

Based on this analysis, the recommendation is to proceed with the acquisition under an active ERM program and three non-negotiable conditions: (1) zero-trust controls must be fully deployed before any live-data integration; (2) a pre-launch sales-engineering team of at least 15 specialists must be in place; and (3) each market must reach an 80 percent driver-adoption rate within six months, monitored monthly with authority to pause expansion if the threshold is not met. Enterprise risk analytics does not eliminate uncertainty, but it converts unquantified exposure into managed probability distributions with clear trigger points—giving decision-makers the structure needed to proceed, monitor, and adjust as conditions evolve.

## REFERENCES

- AI Magazine. (2026, February 23). Amazon utilising AI to automate logistics.
- Amazon.com, Inc. (2025). Annual report on Form 10-K for the fiscal year ended December 31, 2024. U.S. Securities and Exchange Commission.
- Associated Press. (2024, January 29). Amazon's bid to buy Roomba maker iRobot is called off amid pushback in Europe.
- BCG. (2026). Last-mile delivery accounts for 50–60% of total parcel costs. Boston Consulting Group.
- California Attorney General. (2023). California Consumer Privacy Act (CCPA). California Department of Justice.
- CCJ Digital. (2026). Supply chain cyber threats: How malicious AI puts logistics data at risk.
- Committee of Sponsoring Organizations of the Treadway Commission. (2017). Enterprise risk management-Integrating with strategy and performance.
- Cybersecurity and Infrastructure Security Agency. (2019). Binding Operational Directive 19-02: Vulnerability remediation requirements for internet-accessible systems. U.S. Department of Homeland Security. <https://www.cisa.gov/news-events/directives/bod-19-02-vulnerability-remediation-requirements-internet-accessible-systems>
- Deep Tech Nation Switzerland. (2026, March 20). Amazon acquires RIVR: How an ETH Zurich lab built the robot that delivers your packages.
- European Parliament. (2024). Regulation (EU) 2024/1689 (EU Artificial Intelligence Act).
- FedEx Corporation. (2024). FedEx Surround platform overview. FedEx Investor Relations.
- Federal Trade Commission. (2023). FTC v. Amazon.com, Inc.
- Financial Content. (2026, March 19). Amazon (AMZN) 2026 analysis: The \$200 billion AI and aerospace pivot.
- Fleisher, C. S., & Bensoussan, B. E. (2015). Business and competitive analysis (2nd ed.). FT Press.
- HubSpot. (2024). 2024 sales trends report. HubSpot, Inc.
- IBM Security. (2025). Cost of a data breach report 2025. <https://www.ibm.com/reports/data-breach>
- INFORMS. (2016). UPS On-Road Integrated Optimization and Navigation (ORION) - Franz Edelman Award.



International Organization for Standardization. (2018). ISO 31000:2018 Risk management-Guidelines.

McKinsey & Company. (2022). The last-mile delivery challenge.

National Institute of Standards and Technology. (2023). Artificial intelligence risk management framework (AI RMF 1.0).

Oracle Corporation. (2024). Oracle Logistics Cloud platform overview.

RAIN Group. (2023). Win rate benchmarks: Top-performing sales organizations. RAIN Group Center for Sales Research.

Reuters. (2025, October 22). Amazon sees faster delivery speeds with hi-tech driver eyeglasses and AI.

SAP SE. (2024). SAP Transportation Management platform overview.

Shuaibu, A. S., Mahmoud, A. S., & Sheltami, T. R. (2025). A review of last-mile delivery optimization: Strategies, technologies, drone integration, and future trends. *Drones*, 9(3), 158. <https://doi.org/10.3390/drones9030158>

Skadden. (2025, June). M&A in the AI era: Key antitrust and national security considerations.

Statista. (2024). Last-mile delivery share of total shipping costs worldwide, 2018–2023.

Walmart Inc. (2024). Annual report on Form 10-K (AI supply chain transformation, as cited in *AI Magazine*, 2026).

## Appendix 1: Complete Risk Treatment and Response Plan (RT&RP)

The RT&RP below consolidates the three priority risks with qualitative scores (per the canonical register), I&W probabilities, Monte Carlo results, ML AUC, treatment, residual EMV, owner, and KRI triggers. The full 32-risk register and 4-opportunity register (Risk\_HeatMap\_ALY6130\_REVISED.xlsx) is the complete RT&RP and is attached in full.

| <b>PR1 - Commercial Market-Entry / Win-Rate Failure (register cross-ref: R26 / opportunities O2–O3)</b> |                                                                                                                                                                                                                                                                                                                                                                                  |
|---------------------------------------------------------------------------------------------------------|----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| <b>Date Logged</b>                                                                                      | June 3, 2026 (Module 2), updated through Module 6.                                                                                                                                                                                                                                                                                                                               |
| <b>Risk Statement</b>                                                                                   | Failing to achieve expected market adoption of Amazon’s AI logistics platform, caused by entering the commercial logistics-software market without a sales-engineering function, SLA support model, or competitive pricing strategy against vendors with 14+ years of investment, resulting in win-rate collapse below 40%, lost commercial revenue, and eroded acquisition ROI. |
| <b>Causal Chain</b>                                                                                     | No sales-engineering function → incumbents (ORION/Surround/Oracle/SAP) leverage enterprise lock-in → win-rate < 40% → commercial targets missed → ROI erodes.                                                                                                                                                                                                                    |
| <b>Qualitative Score</b>                                                                                | Likelihood 7 × Impact 6 = 42 (High).                                                                                                                                                                                                                                                                                                                                             |
| <b>I&amp;W Probability</b>                                                                              | $P \approx 0.65$ (mean composite warning 0.61, blended 70/30 with register anchor).                                                                                                                                                                                                                                                                                              |
| <b>ML AUC</b>                                                                                           | 0.87 (GBM). Dominant signal: sales-engineering vacancy gap, 2–3 months lead time.                                                                                                                                                                                                                                                                                                |
| <b>3-Point Impact / E[Impact]</b>                                                                       | \$200M / \$500M / \$1,200M; PERT $E[\text{Impact}] = \$567\text{M}$ .                                                                                                                                                                                                                                                                                                            |
| <b>Monte Carlo EMV / P95</b>                                                                            | –\$368M EMV; P95 downside (VaR) = \$890M.                                                                                                                                                                                                                                                                                                                                        |

|                           |                                                                                                                                                                                           |
|---------------------------|-------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| <b>Treatment Strategy</b> | Mitigate: pre-launch competitive benchmarking; hire 15–20 enterprise sales engineers before first RFP; tiered pricing within FTC limits; quarterly win-rate dashboard with VP escalation. |
| <b>Residual EMV</b>       | –\$239M (–35%).                                                                                                                                                                           |
| <b>Risk Owner</b>         | VP, Logistics Technology (primary); Chief Commercial Officer (secondary).                                                                                                                 |
| <b>KRI Red Triggers</b>   | Win-rate < 40% for 2 consecutive quarters; sales-engineering headcount < 75%; feature-gap score < 6.5/10.                                                                                 |
| <b>Review Cadence</b>     | Quarterly win-rate dashboard; monthly sales-engineering headcount review.                                                                                                                 |
| <b>Status</b>             | Open - treatment plan active.                                                                                                                                                             |

|                                                                                                   |                                                                                                                                                                                                                                                                                                                                                                                  |
|---------------------------------------------------------------------------------------------------|----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| <b>PR2 - Cybersecurity Breach of Geolocation and Address Data (register cross-ref: R10 + R11)</b> |                                                                                                                                                                                                                                                                                                                                                                                  |
| <b>Date Logged</b>                                                                                | June 3, 2026 (Module 2), updated through Module 6.                                                                                                                                                                                                                                                                                                                               |
| <b>Risk Statement</b>                                                                             | A cybersecurity breach, ransomware attack, or unauthorized API access against the platform, caused by a materially new attack surface (real-time geolocation, customer addresses, proof-of-delivery photographs) for which existing controls were not designed, resulting in CCPA multi-state notification, enforcement penalties, delivery disruption, and reputational damage. |
| <b>Causal Chain</b>                                                                               | New attack surface + zero-trust not yet deployed → integration-window exposure (months 6–18) → ransomware / unauthorized API access → CCPA 72-hour notification → enforcement, disruption, and reputational damage compound.                                                                                                                                                     |
| <b>Qualitative Score</b>                                                                          | Likelihood 5 × Impact 9 = 45 (High). Asymmetric: moderate probability, extreme tail.                                                                                                                                                                                                                                                                                             |

|                                   |                                                                                                                                                                                                                 |
|-----------------------------------|-----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| <b>I&amp;W Probability</b>        | $P \approx 0.47$ (mean composite 0.44, peaking mid-integration; blended 70/30 with register anchor).                                                                                                            |
| <b>ML AUC</b>                     | 0.89 (GBM). Dominant signal: anomalous geolocation-API traffic, 15–45 days lead time.                                                                                                                           |
| <b>3-Point Impact / E[Impact]</b> | \$250M / \$450M / \$900M; PERT $E[\text{Impact}] = \$492\text{M}$ .                                                                                                                                             |
| <b>Monte Carlo EMV / P95</b>      | –\$231M EMV; P95 downside (VaR) = \$701M.                                                                                                                                                                       |
| <b>Treatment Strategy</b>         | Mitigate + Transfer: zero-trust as a hard launch gate; API security gateway with anomaly detection; logistics-specific cyber insurance (CCPA + business interruption up to \$50M/incident); tested IR playbook. |
| <b>Residual EMV</b>               | –\$104M (–55%) - the largest proportional reduction in the portfolio.                                                                                                                                           |
| <b>Risk Owner</b>                 | CISO / VP Information Security (primary); General Counsel for CCPA (secondary).                                                                                                                                 |
| <b>KRI Red Triggers</b>           | Any confirmed unauthorized geo/address API access; any critical vuln > 15 days; zero-trust coverage < 90%.                                                                                                      |
| <b>Review Cadence</b>             | Continuous geo-API monitoring; weekly CISO vulnerability review; monthly IAM coverage audit.                                                                                                                    |
| <b>Status</b>                     | Open - zero-trust deployment is the highest-priority pre-launch action.                                                                                                                                         |

|                                                                                                              |                                                                                                                                                                                                                            |
|--------------------------------------------------------------------------------------------------------------|----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| <b>PR3 - Cost-Reduction and Sustainability Leadership (Opportunity) (register cross-ref: Opportunity O1)</b> |                                                                                                                                                                                                                            |
| <b>Date Logged</b>                                                                                           | June 3, 2026 (Module 2), updated through Module 6.                                                                                                                                                                         |
| <b>Opportunity Statement</b>                                                                                 | Significant cost reduction and sustainability leadership enabled by AI route optimization, resulting in 15–25% per-package cost reduction, measurable CO <sub>2</sub> decreases, and positioning ahead of California Clean |

|                                   |                                                                                                                                                                                                                                                  |
|-----------------------------------|--------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
|                                   | Miles and EU urban-access mandates - generating \$400M–\$1.36B in annual savings and turning compliance into a competitive moat.                                                                                                                 |
| <b>Causal Chain</b>               | AI optimization across 8M daily packages → 15–25% per-package reduction → \$400M–\$1.36B annual savings + CO <sub>2</sub> reduction → pre-compliance with Clean Miles / EU access → urban delivery-access advantage over non-compliant carriers. |
| <b>Qualitative Score</b>          | Likelihood 9 × Impact 8 = 72 (High upside).                                                                                                                                                                                                      |
| <b>I&amp;W Probability</b>        | $P \approx 0.62$ ( $P(\text{realized}) = 1 - \text{mean warning } 0.44$ , blended 70/30 with register anchor).                                                                                                                                   |
| <b>ML AUC</b>                     | 1.00 (Logistic, synthetic data). Dominant signal: driver adoption rate.                                                                                                                                                                          |
| <b>3-Point Impact / E[Impact]</b> | +\$400M / +\$700M / +\$1,200M; PERT E[Impact] = +\$733M.                                                                                                                                                                                         |
| <b>Monte Carlo EMV / P95</b>      | +\$453M EMV; P95 best case = +\$991M.                                                                                                                                                                                                            |
| <b>Treatment Strategy</b>         | Exploit + Enhance: accelerate rollout from 3-city pilot to 50+ stations in 18 months; structured driver-adoption program (80% target per market); CA CARB Clean Miles pre-compliance filing; public CO <sub>2</sub> KPI reporting.               |
| <b>Residual EMV</b>               | +\$534M (+18%).                                                                                                                                                                                                                                  |
| <b>Risk Owner</b>                 | COO, Last-Mile Delivery (primary); Chief Sustainability Officer (secondary).                                                                                                                                                                     |
| <b>KRI Red Triggers</b>           | Per-package reduction < 5%; driver adoption < 60%; OTDPI < 97 for 2 consecutive months.                                                                                                                                                          |
| <b>Review Cadence</b>             | Monthly per-package cost; monthly driver adoption by market; quarterly sustainability reporting.                                                                                                                                                 |
| <b>Status</b>                     | Open - pilot expansion active; adoption program in design.                                                                                                                                                                                       |

|                                                                                                                                                                            |
|----------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| <b>Portfolio Summary - Net Expected Value After Treatment</b>                                                                                                              |
| <b>PR1 residual –\$239M   PR2 residual –\$104M   PR3 enhanced +\$534M   NET +\$191M positive expected value, before the strategic option value of the platform itself.</b> |

*Note on residual figures: the residual-EMV reductions (–35% / –55% / +18%) are modeled estimates of treatment efficacy recorded in the team’s Risk\_Mitigation deliverable; they are management judgments of how much each treatment lowers exposure, not direct outputs of the Monte Carlo notebook.*